

Ex1a: Text Generation using N-gram Language model

Learning Objective:

Implement an N-Gram-based predictive text model using bigram and trigram probability distributions. This helps to understand how contextual probability influences next-word prediction.

Steps:

1. Initialize the environment by installing and importing the Natural Language Toolkit (nltk). Download the **Brown Corpus** (for text data) and the **Universal Tagset** (for simplified Part-of-Speech tags).
2. Extract raw sentences from the Brown Corpus using brown.sents().
3. **Tokenization:** Convert all words to lowercase and flatten the nested list into a single list of tokens to verify the total volume of training data.
4. Build an N-Gram Model for Bigram and Trigram.
5. Predict Next Word and display Top-5 Suggestions
6. Test Bigram & Trigram Prediction on certain words.

Program:

```
import nltk
from nltk.corpus import brown
from collections import defaultdict, Counter
import random
from collections import defaultdict, Counter
import nltk
from nltk.corpus import brown
nltk.download('brown')
words = [word.lower() for word in brown.words() if word.isalpha()]
def build_ngram_model(words, n=3):
    ngram_counts = defaultdict(Counter)
    for i in range(len(words) - n + 1):
        context = tuple(words[i:i+n-1])
        target = words[i+n-1]
        ngram_counts[context][target] += 1
    return ngram_counts

def show_next_word_counts(context, model):
    context = tuple(context.lower().split())
    if context not in model:
        print("No prediction available")
    return
```

```
total_count = sum(model[context].values())
print(f"Context: {''.join(context)}")
print(f"TOTAL count: {total_count}\n")
print("Next word counts:\n")
for word, count in model[context].most_common():
    print(f'{word}<15} : {count}')
# Example usage
N = 2
ngram_model = build_ngram_model(words, N)
context = "Apple"
show_next_word_counts(context, ngram_model)
```

Output:

Context: apple

TOTAL count: 9

Next word

counts:

pie : 2

tree : 2

trees : 1

pies : 1

cider : 1

the : 1

and : 1

Learning Outcome:

Upon completion of this experiment, the bigram and trigram probability distributions for next-word prediction was implemented and evaluated using the Brown Corpus.